

MapReduce & Hadoop

- Juanjo Mostazo
- SeedRocket BCN
- June 2011




M/R: Motivation

- Process big amount of data to produce other data
- Scale up VS Scale out



Scale Oracle Real-Time Collaboration Depending on Your Requirements



M/R: What is it?

- Different programming paradigm
- Based on a google paper (2004)
- Automatic parallelization and distribution
- I / O Scheduling
- Fault tolerance
- Status and monitoring



M/R: Basics

- Input & Output: set of key / value pairs
- Big amount of data group & sort
- Job = Two phases = Mapper & Reducer
- Map $(in_key, in_value) \rightarrow$
list $(interm_key, interm_value)$
- Reduce $(interm_key, list(interm_value)) \rightarrow$
list (out_key, out_value)

```

map(String input_key,
      String input_value):
  for each shape s in input_value:
    EmitIntermediate(s, 1);

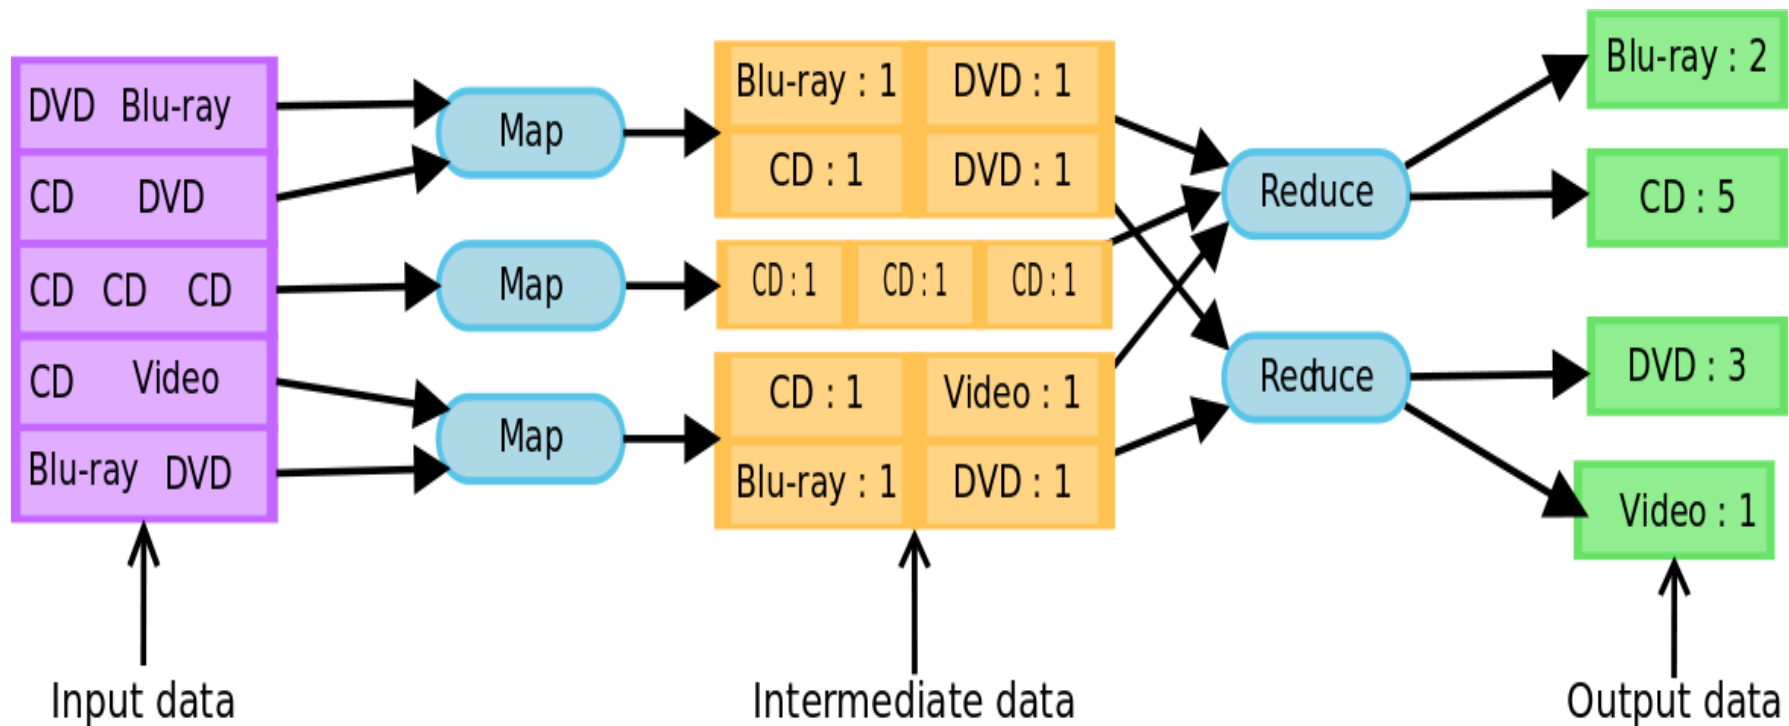
```

```

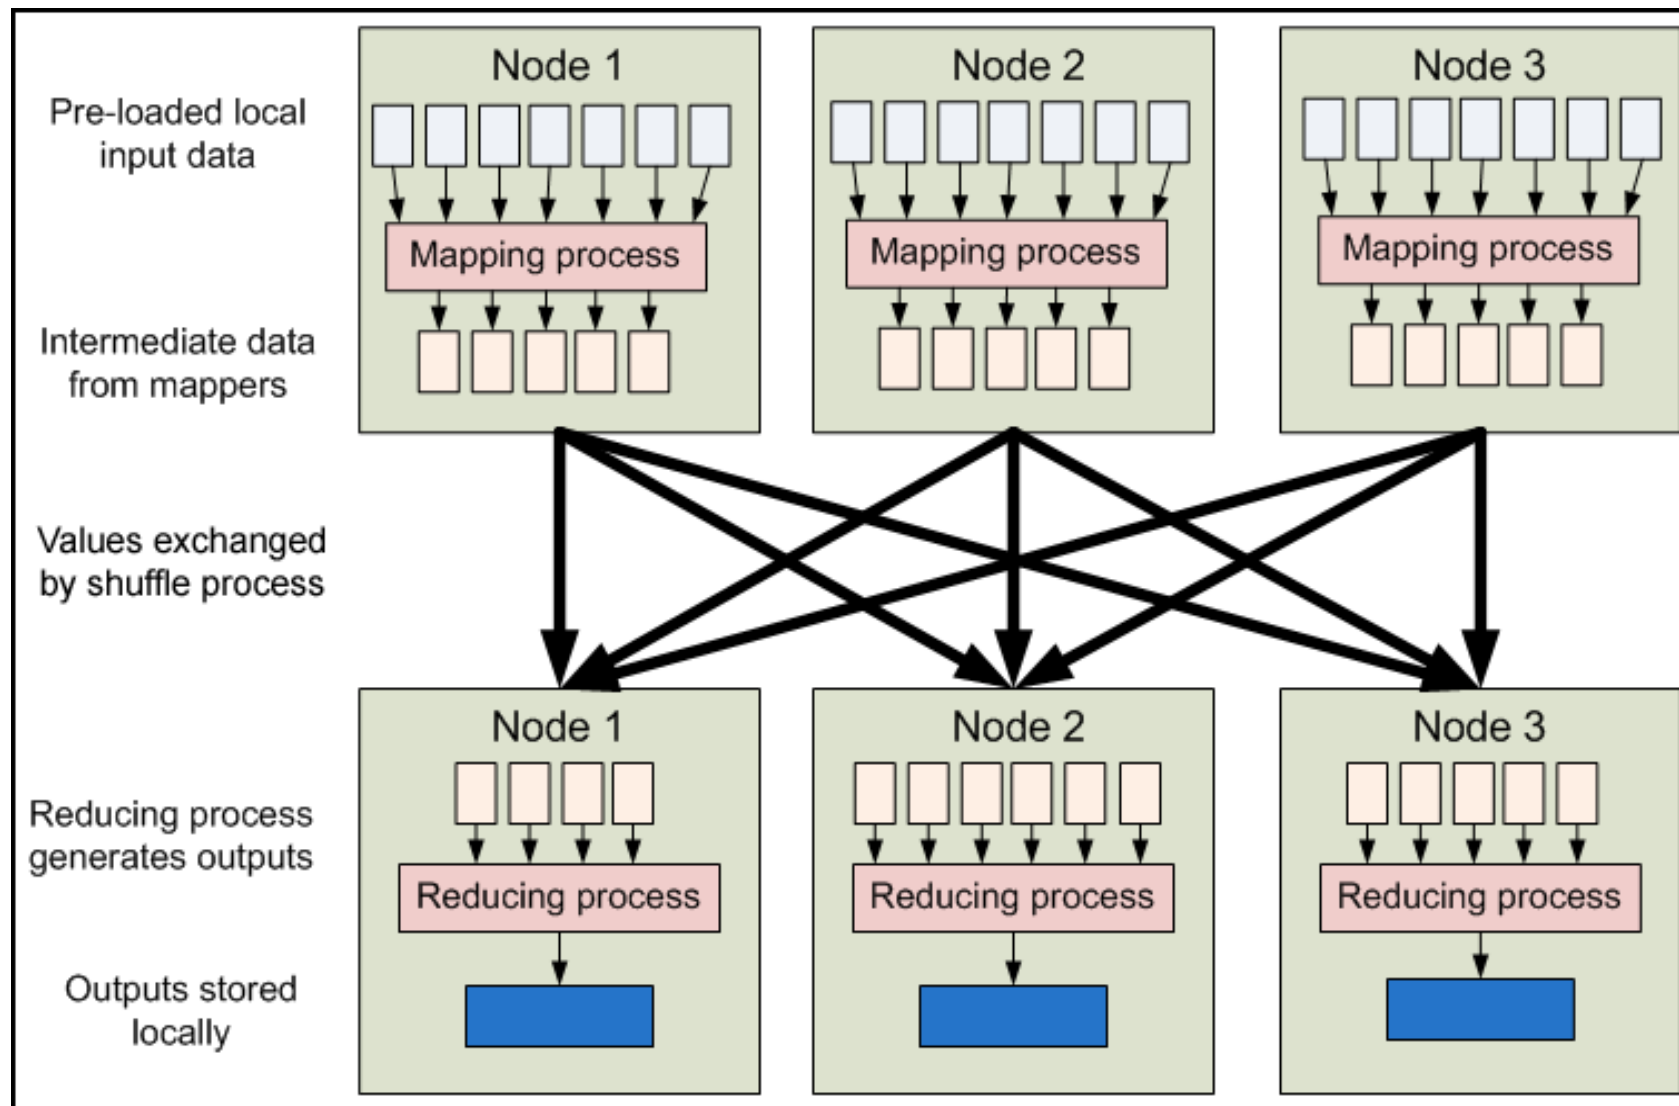
reduce(String interm_key,
         Iterator interm_values):
  result = 0;
  for each v in interm_values:
    result += v;
  Emit(interm_key, result);

```

M/R: Example (word counter)



M/R: Workflow



Hadoop: What is it?

- Framework based on GMR / GFS
- Apache project
- Developed in Java
- Multiple applications
- Used by many companies
- And growing!



facebook

Linked in


Hadoop: HDFS dist. systems

- Data is sent to computation
(whereas here computation goes to data),
- Nodes must get info from each other
- Very difficult to recover from partial failure

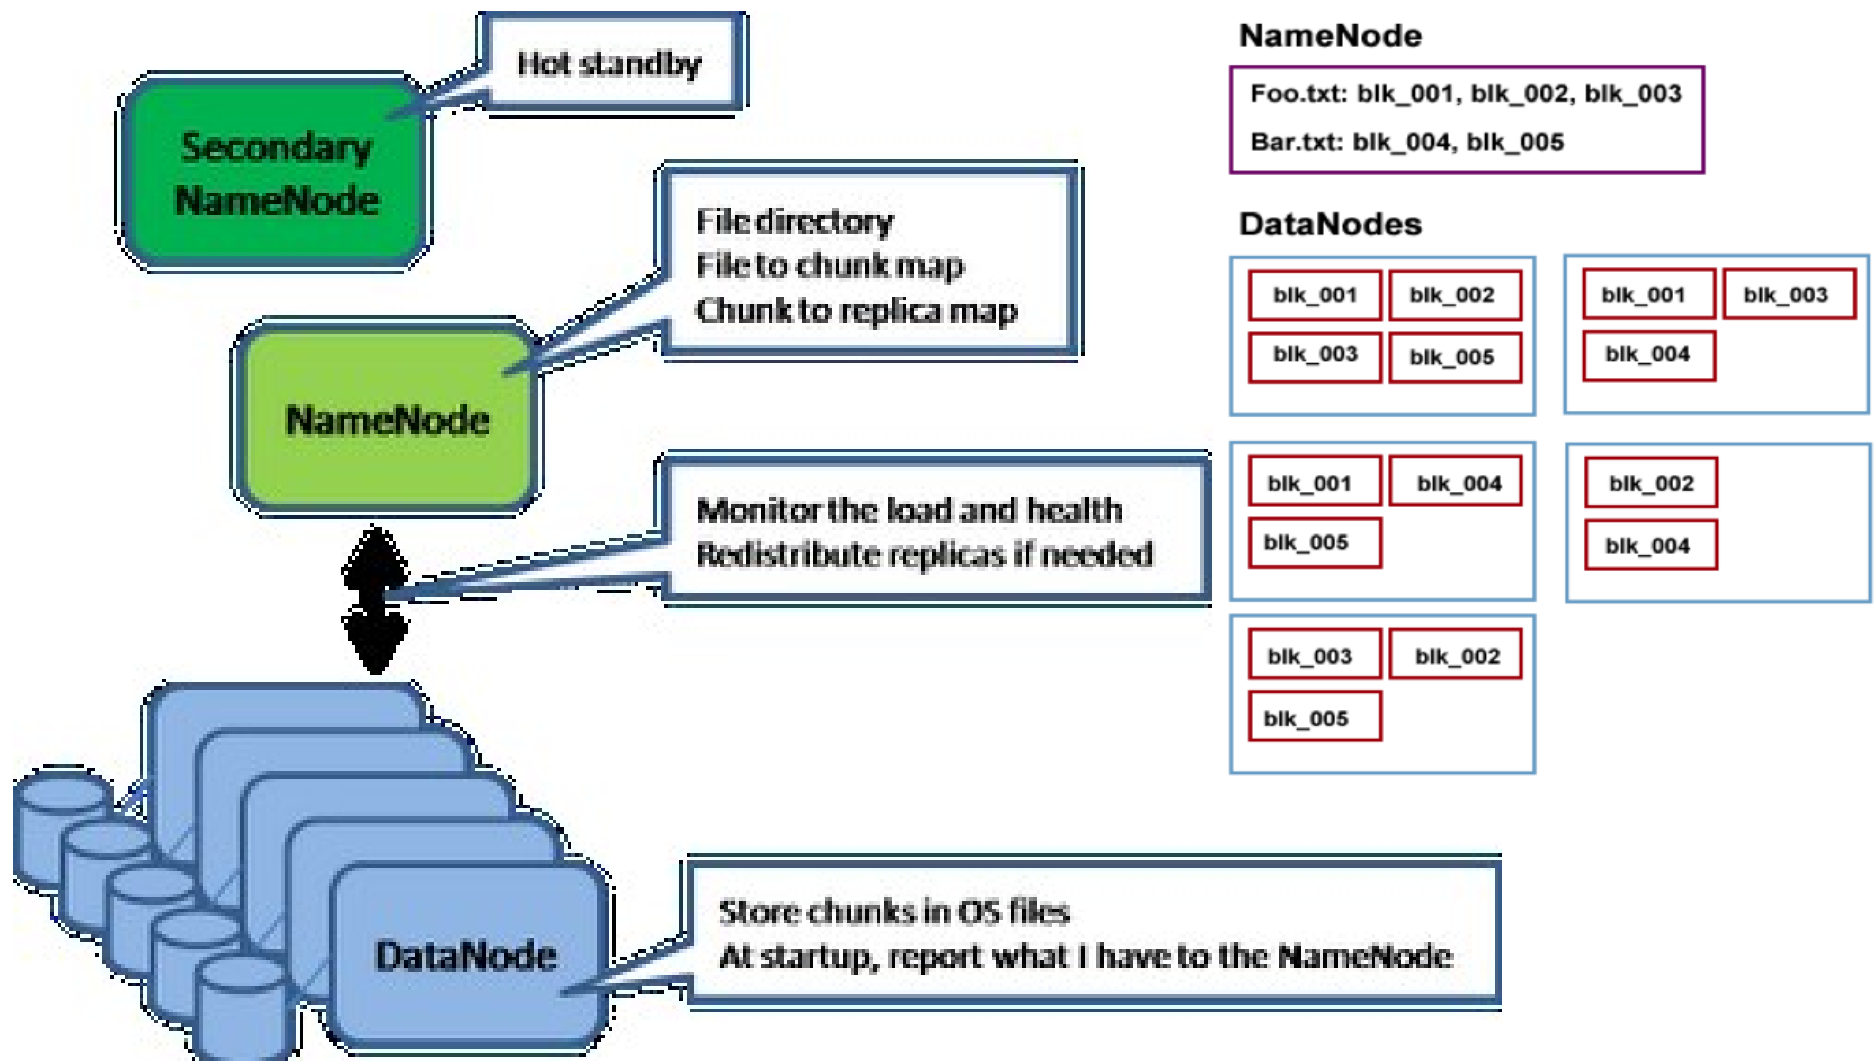


Hadoop: HDFS concepts


- Distributed file system.
Layer on top ext3, xfs...
- Works better on huge files
- Redundancy (default 3)
- Bad seeking, no append!
- Good rack scale. Not good data center scale
- File divided in 64Mb – 128Mb blocks



Hadoop: HDFS Architecture




Hadoop: Architecture v1

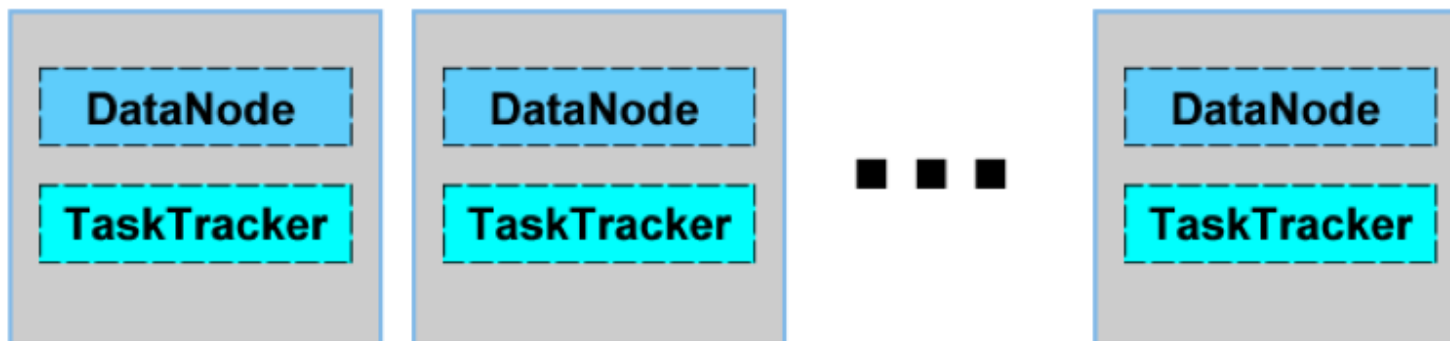


Hadoop: Architecture v2

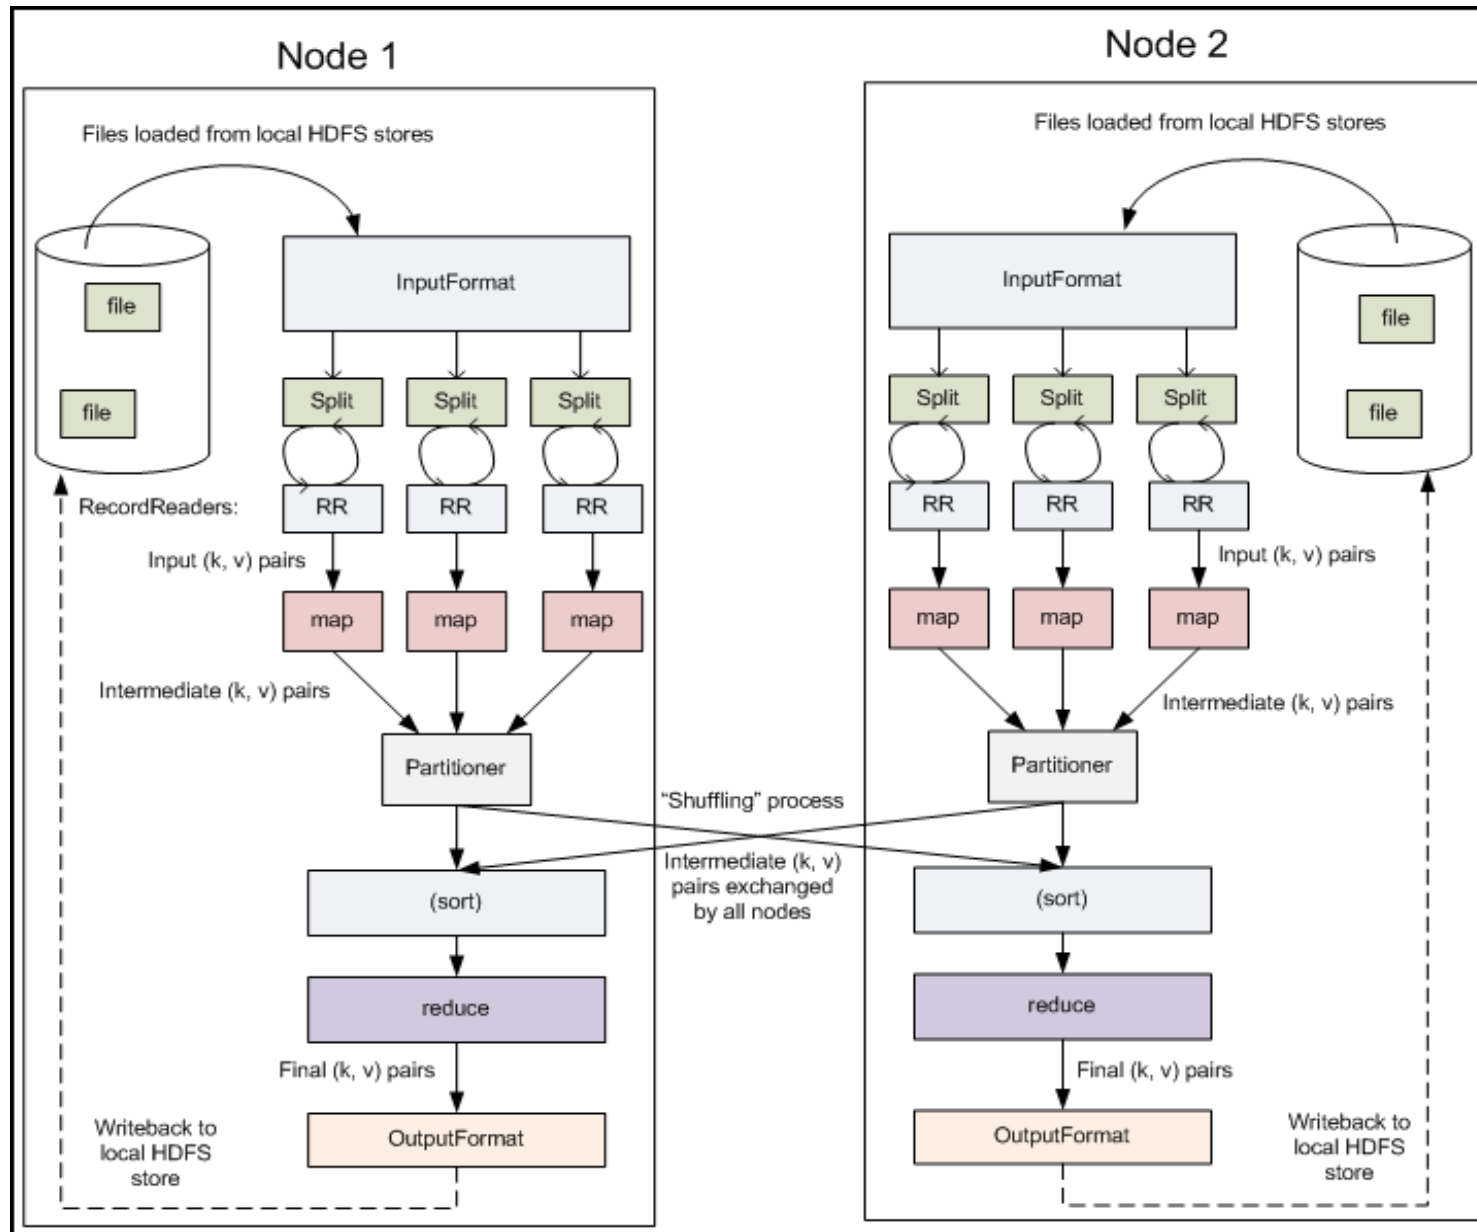
Master Nodes



Slave Nodes



Hadoop: Architecture v3



```

map(String input_key,
      String input_value):
  for each shape s in input_value:
    EmitIntermediate(s, 1);


```

```

reduce(String interm_key,
         Iterator interm_values):
  result = 0;
  for each v in interm_values:
    result += v;
  Emit(interm_key, result);

```


M/R: Example (word counter)



Hadoop: EcoSystem



• FLUME



• ZOOKEEPER




Hadoop: Advanced stuff

- Distributed caches
- Partitioner
- Sort comparator
- Group comparator
- Combiner
- Input format & Record reader
- MultiInput
- MultiOutput
- Compression (LZO)



Hadoop: Conclussions

- Simplify large-scale computation
- Hide parallel programming issues
- Easy to get into & develop
- Deeply used & maintained by community
- Possibility yo throw away RDBMs! (Bottleneck)



MapReduce & Hadoop

THANKS!

- Juanjo Mostazo
- juanj.mostazo@gmail.com

